

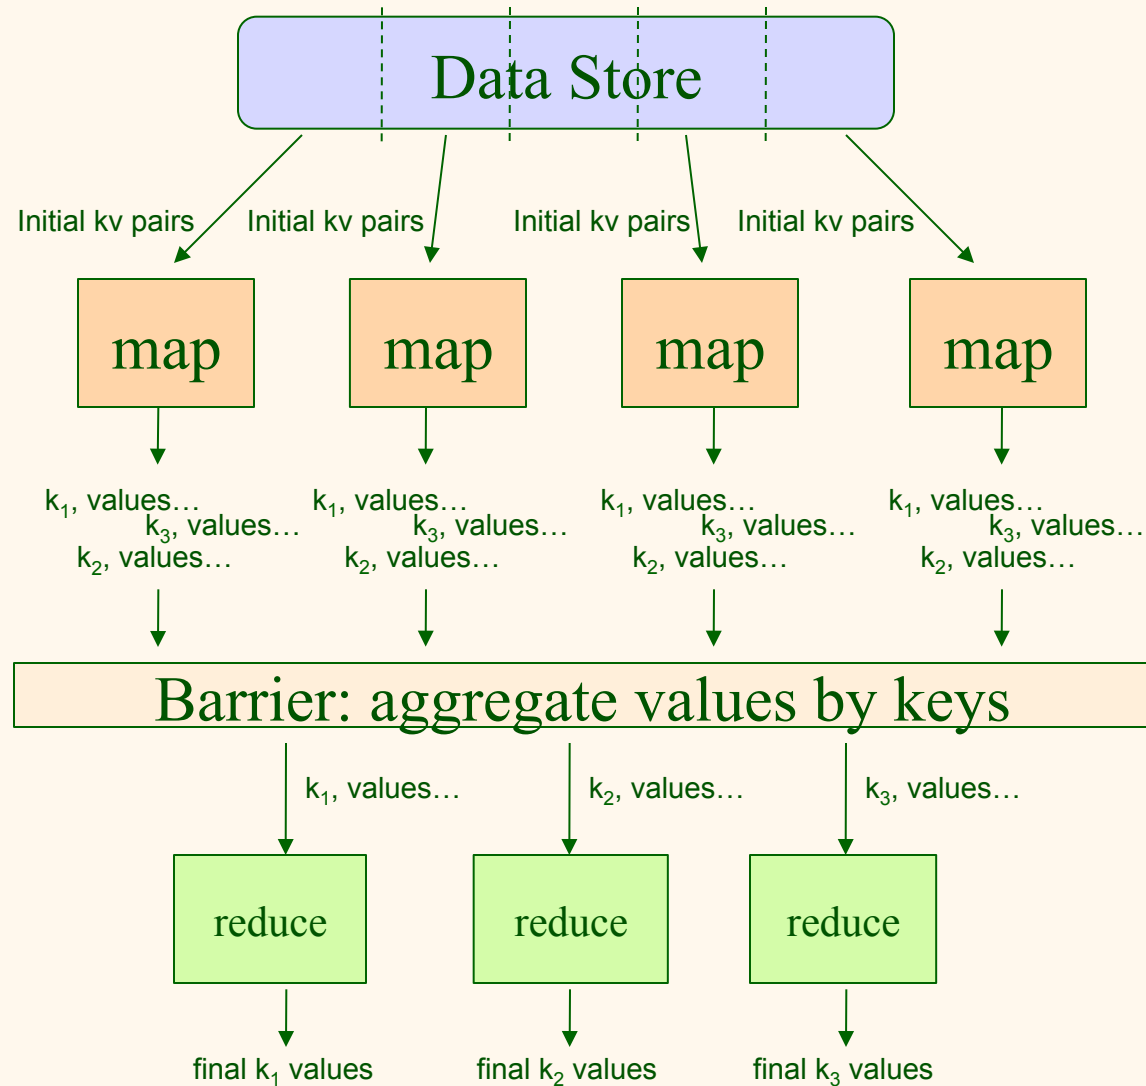
Map Reduce Continued



Example: Word Count


```
map(String key, String value):  
    // key: document id  
    // value: document contents  
    for each word w in value:  
        EmitIntermediate(w, "1");  
  
reduce(String key, Iterator values):  
    // key: a word  
    // values: a list of counts  
    int result = 0;  
    for each v in values:  
        result += ParseInt(v);  
    Emit(key, AsString(result));
```

Map Reduce



SSSP - Dijkstra's Algorithm

Single-Source-Shortest-Path





SSSP in Map Reduce


- ❖ Can't run Dijkstra's algorithm directly
 - Can't have a global queue!
- ❖ Another way to do it: Parallel BFS
- ❖ Start by assuming all edge weights are equal
 - Will relax this later




Finding the Shortest Path

- ❖ Intuition: process all nodes at each step
- ❖ Some nodes have no information (distance = infinity)
 - So can't do much
- ❖ But other nodes do know something
 - E.g. source knows it is at distance 0
 - So can pass this fact on to its out-neighbors
 - ◆ Who now know they are at distance 1!
 - At the next iteration, these neighbors know they're at distance 1
 - ◆ So can tell *their* out-neighbors they're at distance 2.

Parallel BFS



From Intuition to Algorithm



Map	Reduce
(B,1)	(B,[1])
(C,1)	(C,[1])
<hr/>	
(D,2)	(B,[1,2])
(B,2)	(C,[1])
(B,1)	(D,[1])
(C,1)	(D,[2])

- ❖ A map task receives
 - Key: node n
 - Value: D (distance from start), points-to (list of nodes reachable from n)
- ❖ $\forall p \in \text{points-to: emit } (p, D+1)$
- ❖ The reduce task gathers possible distances to a given p and selects the minimum one
- ❖ Possible through the magic of the "sort and shuffle"
between Map and Reduce
 - Map processes node and updates distances of **out-neighbors**
 - Reduce processes node based on info from its **in-neighbors**



Multiple Iterations Needed

- ❖ Each Map Reduce task advances the “known frontier” by one hop
 - Subsequent iterations include more reachable nodes as frontier advances
 - Multiple iterations are needed to explore entire graph
 - Feed output back into the same MapReduce task



Multiple Iterations Needed

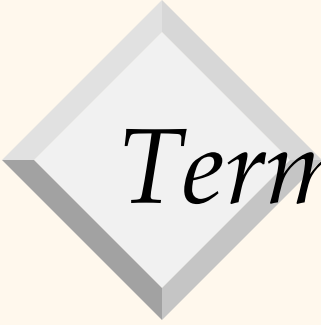
- ❖ Passing along the graph structure:
 - Next iteration of Map needs points-to list again
 - So need to "carry" it with us as we run the algorithm



```
class MAPPER
  method MAP(nid  $n$ , node  $N$ )
     $d \leftarrow N.DISTANCE$ 
    EMIT(nid  $n$ ,  $N$ )
    for all nodeid  $m \in N.ADJACENCYLIST$  do
      EMIT(nid  $m$ ,  $d + 1$ )
```

Reduce

```
class REDUCER
  method REDUCE(nid  $m$ , [ $d_1, d_2, \dots$ ])
     $d_{min} \leftarrow \infty$ 
     $M \leftarrow \emptyset$ 
    for all  $d \in \text{counts } [d_1, d_2, \dots]$  do
      if IsNode( $d$ ) then
         $M \leftarrow d$ 
      else if  $d < d_{min}$  then
         $d_{min} \leftarrow d$ 
     $M.\text{DISTANCE} \leftarrow d_{min}$ 
    EMIT(nid  $m$ , node  $M$ )
```



Termination


- ❖ Eventually, all nodes will be discovered, all edges will be considered (in a connected graph)
- ❖ Stop when there are no nodes with a distance of infinity
 - Can be checked by the driver/harness/program that runs the outer loop and schedules each Map Reduce job




Weighted Edges

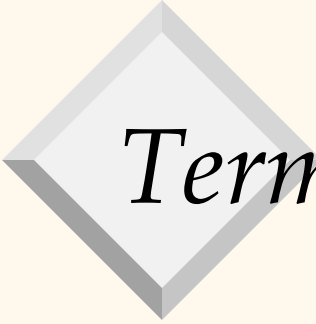
- ❖ Now add positive weights to the edges
- ❖ Simple change: points-to list in map task includes a weight w for each pointed-to node
 - emit $(p, D+w_p)$ instead of $(p, D+1)$ for each node p
- ❖ Termination behavior different
 - Just because we've reached a node doesn't mean we've found the shortest path to it!

Node Exploration Process



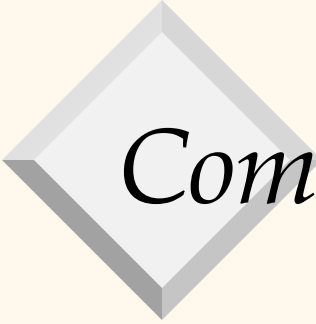
Node Exploration Process





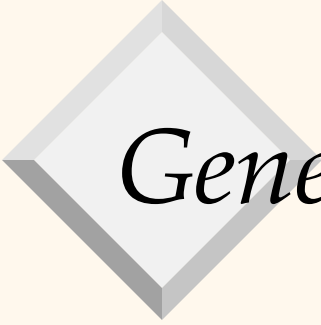
Termination

- ❖ When distances have not changed during an iteration, safe to stop



Comparison to Dijkstra

- ❖ Dijkstra's algorithm is more efficient
 - At any step it only pursues edges from the minimum-cost path inside the frontier
 - Only processes each node once
- ❖ MapReduce explores all paths in parallel
 - Does a lot of recomputation
 - ◆ Not a bug, need it to handle situations where the "shortest" path contains more edges than another available path
 - But can be done in parallel



General Approach

- ❖ Graph algorithms with MapReduce:
 - Each map task receives a node and its outlinks
 - Map task compute some function of the link structure, emits value with target as the key
 - Reduce task collects keys (target nodes) and aggregates
- ❖ Iterate multiple MapReduce cycles until some termination condition

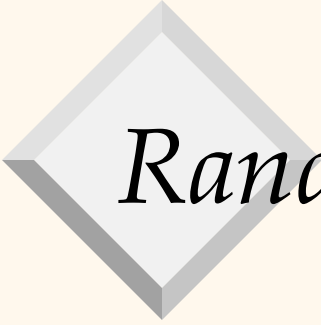


PageRank

- ❖ Google's famous algorithm for ranking Web Pages
 - A measure of "quality reputation" of a page
 - Useful for ranking/ordering search results


PageRank

- ❖ Based on link structure (graph) of the web
- ❖ Intuition from academic citation network:
 - Lots of citations (incoming links) = probably a high quality paper
 - If a high quality paper cites your paper, your paper is probably of high quality too
 - Vice versa not necessarily true (just by citing a high quality paper you don't make your paper high quality 😊)



Random Surfer Model

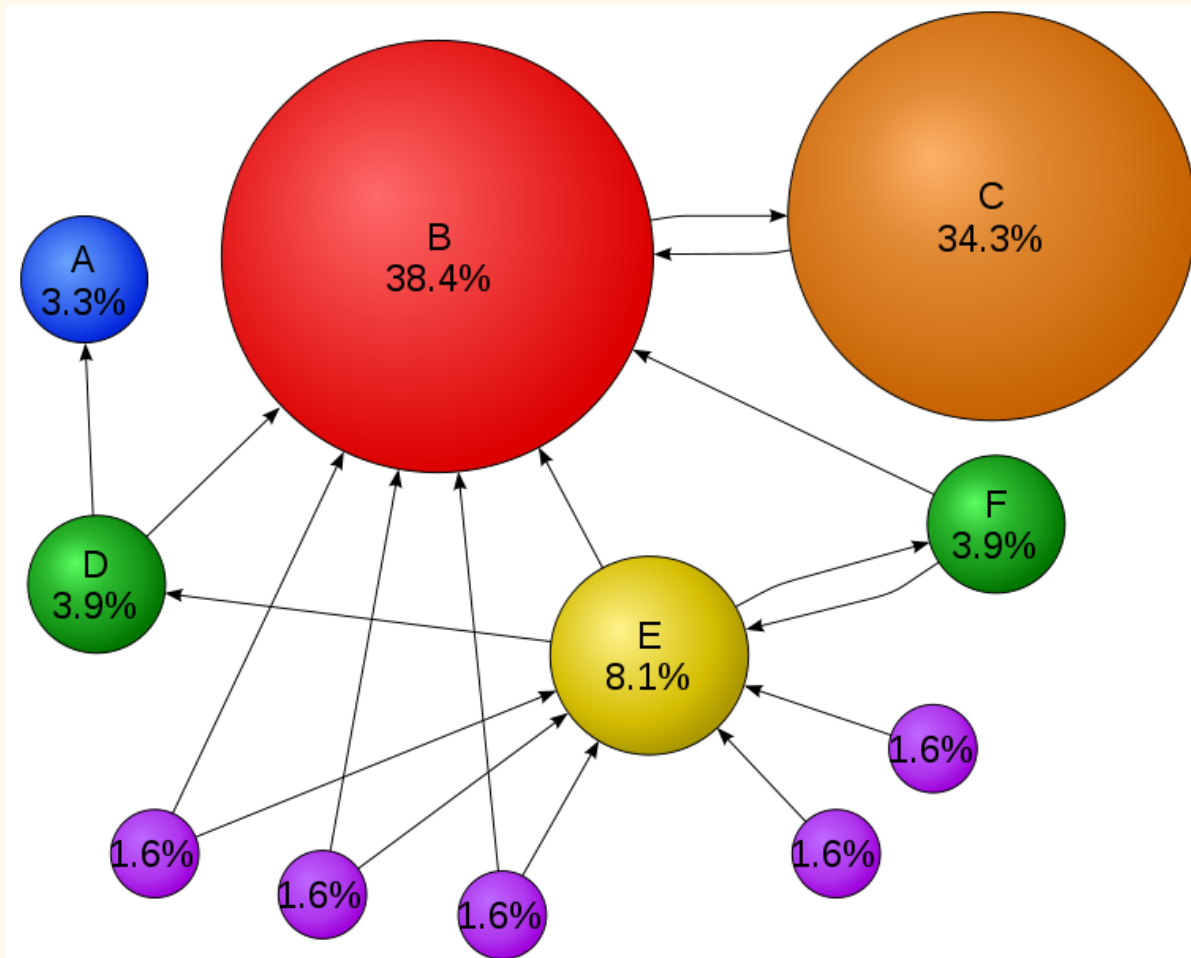
- ❖ Another intuition for PageRank
- ❖ Imagine a surfer who starts on a randomly chosen page and then follows outgoing links at random
 - Markov process
- ❖ PageRank is probability that user will arrive at a given page during this random walk

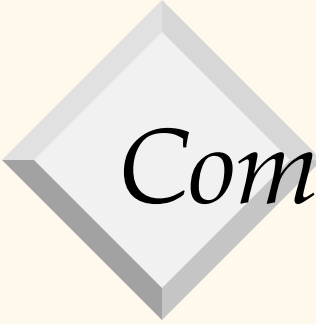


A little more complex!

- ❖ Model assumes that surfer doesn't always follow a link, but sometimes e.g. bookmarks instead.
- ❖ Before each move, surfer flips a coin
 - With probability $1 - \alpha$, follows an out-link
 - With probability α , teleports to a (uniformly chosen) random page

PageRank Example







Computing PageRank

- ❖ Let's start with some simplifying assumptions
 - Assume all nodes have at least one outlink
 - And surfer never does a random restart using bookmarks
- ❖ Will talk about how to lift these assumptions soon

Simplified PageRank Intuition

- ❖ PageRank of a page is based on the PageRank of the pages which link to it
- ❖ A page divides its PageRank equally among all its outgoing links

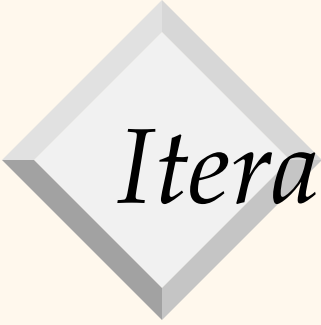




Somewhat more formally

$$P(n) = \sum_{m \in L(n)} \frac{P(m)}{C(m)}$$

- ❖ $L(n)$ is the set of pages that link to n and $C(m)$ is the number of out-neighbors of page m





Iterative computation of SPR

- ❖ How do we compute this?
- ❖ Various methods, but we are interested in Map Reduce
- ❖ Idea:
 - Initialize everything to the same PageRank ($1/\text{number of nodes}$)
 - "pass around" PageRank contributions from nodes to their out-neighbors



Example

Iteration 1



Example

Iteration 2





```
class MAPPER
```

```
  method MAP(nid  $n$ , node  $N$ )
```

```
     $p \leftarrow N.\text{PAGERANK} / |N.\text{ADJACENCYLIST}|$ 
```

```
    EMIT(nid  $n$ ,  $N$ ) ▷ Pass along graph structure
```

```
    for all nodeid  $m \in N.\text{ADJACENCYLIST}$  do
```

```
      EMIT(nid  $m$ ,  $p$ ) ▷ Pass PageRank mass to neighbors
```

Reduce

```
class REDUCER
  method REDUCE(nid  $m$ ,  $[p_1, p_2, \dots]$ )
     $M \leftarrow \emptyset$ 
    for all  $p \in \text{counts } [p_1, p_2, \dots]$  do
      if ISNODE( $p$ ) then
         $M \leftarrow p$  ▷ Recover graph structure
      else
         $s \leftarrow s + p$  ▷ Sum incoming PageRank contributions
     $M.\text{PAGERANK} \leftarrow s$ 
    EMIT(nid  $m$ , node  $M$ )
```




Iterate

- ❖ Full algorithm is iterative
- ❖ Initialize the nodes to uniform distribution
- ❖ Run the two MR jobs described iteratively
- ❖ Until convergence (no change)



Now, back to dangling nodes

- ❖ If any PageRank is lost due to nodes with no out-neighbors, redistribute that PageRank uniformly throughout the graph for next iteration
- ❖ In the Map Reduce model: keep track of any lost PageRank
 - E.g. by using a special reserved intermediate key, or using a counter (i.e. storing it somewhere)



Solution

- ❖ After the MR task is done, do a cleanup pass
- ❖ Deal both with "missing mass" and with random restart factor



Cleanup pass

m:missing
G:number of all nodes

- ❖ Adjust the PageRank of each node to be

$$p' = \alpha \left(\frac{1}{|G|} \right) + (1 - \alpha) \left(\frac{m}{|G|} + p \right)$$

- ❖ Where p is the current PageRank, m is the mass lost due to sinks, and $|G|$ is the number of nodes in the graph
- ❖ This can be done using a map job (no reduce)



The full PageRank Algorithm

- ❖ Initialize the nodes to uniform distribution
- ❖ Run the two MR jobs described iteratively
- ❖ Until convergence (no change)



You can now start H4

- ❖ Available in CMS
- ❖ PageRank in Hadoop
 - And a second part on Neo4j (covered after break)
- ❖ Installation highly nontrivial
 - Consider starting over Spring Break if you are not experienced with command line Linux (or OS X)